

Multicompetitor and ranking models

Goal for this set of lecture notes

- Normal models for measured outcomes in multicompetitor events
- Plackett-Luce model for rank orderings
- Applications



Goal for this set of lecture notes

- The normal models are a special case of blocked one-way ANOVA. One reference (not specific to sports!) is
 - Section 9.2 of Albert, J., Rizzo, M. (2012). *Analysis of Variance II*. In: *R by Example. Use R!*. Springer, New York, NY. (this chapter is uploaded to the course Canvas site)
- Here are two references for the Plackett-Luce model:
 - Turner, H.L., van Etten, J., Firth, D. *et al.* (2020) Modelling rankings in R: the **PlackettLuce** package. *Comput Stat* **35**, 1027–1057.
 - The start of chapter 3 of Zhang, S. (2021). *Building upon Bradley-Terry and Plackett-Luce: some methods for modeling paired comparison and rank order data* (Doctoral dissertation, Harvard University).

Introduction

Interest in measuring ability or merit in multicompetitor games/sports in which competitors receive numerical scores, or are rank-ordered, such as

- races (human, car, horses)
- poker
- gymnastics
- diving
- golf

Introduction

We will focus on two types of outcomes:

- Multicompetitor games/sports in which outcomes are numerical (e.g., race times, point scores, etc.)
- Multicompetitor games/sports which result only in rank orderings

Even when outcomes are numerical, we can choose to use rank orderings and ignore the numerical values.

This might make sense in scenarios when we want a more robust analysis, particularly to the distributional assumption for the numerical scores.

Quantitative outcome models

The situation:

- Suppose we have J competitors, a subset of whom compete in n multiplayer games.
- For example, suppose we observe J athletes who compete in n races, and that we measure the finish time for each race.
- For game $i = 1, \dots, n$, let y_{ij} be the score of player j in game i assuming the player competed.

Quantitative outcome models

Assuming that the y_{ij} are roughly normally distributed within game i , then the model we can assume is

$$y_{ij} = \beta_i + \theta_j + \varepsilon_{ij}$$

Contribution of game i

Sampling error

where, independently across players and games,

$$\varepsilon_{ij} \sim N(0, \sigma^2)$$

Thus, the model assumes $E(y_{ij}) = \beta_i + \theta_j$

Quantitative outcome models

- This model assumes that the strength of competitor j is represented by θ_j .
- Also, the model assumes that game/race conditions may be distinct from each other, e.g., due to temperature, humidity, wind, or other environmental conditions.
- Thus the model controls for the effect of game i by including β_i as an additive term.

Problem: As specified, the model is unidentifiable! Adding the same value to all the θ_j and subtracting that value to all the β_i results in the same mean for y_{ij} .

This seems to be a common theme in these linear models!

Quantitative outcome models

Two solutions:

1. Impose a linear constraint on the θ_j . This can be done by requiring that $\sum_{j=1}^J \theta_j = 0$.
2. Perform ridge regularization on the θ_j (but not on the β_i).

We will examine each solution briefly, since we have seen them in other contexts already.

Quantitative outcome models

Suppose four players, A, B, C and D , compete in two races.

The results were:

Race 1: $y_{1B} = 4.1, y_{1C} = 3.9, y_{1A} = 2.8, y_{1D} = 2.6$

Race 2: $y_{1B} = 3.8, y_{1A} = 3.5, y_{2C} = 3.3, y_{2D} = 2.7$

We now create two X matrices encoding the players and the races, each with 8 rows ($= 4 \text{ players} \times 2 \text{ races}$).

Quantitative outcome models

Race 1: $y_{1B} = 4.1, y_{1C} = 3.9, y_{1A} = 2.8, y_{1D} = 2.6$

Race 2: $y_{1B} = 3.8, y_{1A} = 3.5, y_{2C} = 3.3, y_{2D} = 2.7$

$$\begin{array}{c}
 \mathbf{X}_\theta = \begin{bmatrix}
 \textcolor{blue}{A} & \textcolor{blue}{B} & \textcolor{blue}{C} & \textcolor{blue}{D} \\
 0 & 1 & 0 & 0 \\
 0 & 0 & 1 & 0 \\
 1 & 0 & 0 & 0 \\
 0 & 0 & 0 & 1 \\
 0 & 1 & 0 & 0 \\
 1 & 0 & 0 & 0 \\
 0 & 0 & 1 & 0 \\
 0 & 0 & 0 & 1
 \end{bmatrix}
 \begin{array}{l}
 \textcolor{blue}{\text{Race 1, place 1}} \\
 \textcolor{blue}{\text{Race 1, place 2}} \\
 \textcolor{blue}{\text{Race 1, place 3}} \\
 \textcolor{blue}{\text{Race 1, place 4}} \\
 \textcolor{blue}{\text{Race 2, place 1}} \\
 \textcolor{blue}{\text{Race 2, place 2}} \\
 \textcolor{blue}{\text{Race 2, place 3}} \\
 \textcolor{blue}{\text{Race 2, place 4}}
 \end{array}
 \end{array}
 \quad
 \mathbf{X}_\beta = \begin{array}{c}
 \begin{bmatrix}
 \textcolor{blue}{R1} & \textcolor{blue}{R2} \\
 1 & 0 \\
 1 & 0 \\
 1 & 0 \\
 1 & 0 \\
 0 & 1 \\
 0 & 1 \\
 0 & 1 \\
 0 & 1
 \end{bmatrix}
 \begin{array}{l}
 \textcolor{blue}{\text{Race 1, place 1}} \\
 \textcolor{blue}{\text{Race 1, place 2}} \\
 \textcolor{blue}{\text{Race 1, place 3}} \\
 \textcolor{blue}{\text{Race 1, place 4}} \\
 \textcolor{blue}{\text{Race 2, place 1}} \\
 \textcolor{blue}{\text{Race 2, place 2}} \\
 \textcolor{blue}{\text{Race 2, place 3}} \\
 \textcolor{blue}{\text{Race 2, place 4}}
 \end{array}
 \end{array}
 \quad
 \mathbf{y} = \begin{bmatrix}
 4.1 \\
 3.9 \\
 2.8 \\
 2.6 \\
 3.8 \\
 3.5 \\
 3.3 \\
 2.7
 \end{bmatrix}$$

Quantitative outcome models

Impose a linear constraint on the θ_j through $\sum_{j=1}^J \theta_j = 0$.

Let

$$W = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ -1 & -1 & -1 \end{bmatrix}$$

$$X_{\theta}^* = X_{\theta} W = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ -1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \\ -1 & -1 & -1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \\ -1 & -1 & -1 \end{bmatrix}$$

Quantitative outcome models

Impose a linear constraint on the θ_j through $\sum_{j=1}^J \theta_j = 0$.

Now we can fit the model using X_{θ}^* and X_{β} as the predictors.

$$E(\mathbf{y}) = X_{\theta}^* \boldsymbol{\theta}_{-1} + X_{\beta} \boldsymbol{\beta} =$$
$$\begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \\ -1 & -1 & -1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \\ -1 & -1 & -1 \end{bmatrix} \begin{bmatrix} \theta_A \\ \theta_B \\ \theta_C \end{bmatrix} + \begin{bmatrix} 1 & 0 \\ 1 & 0 \\ 1 & 0 \\ 1 & 0 \\ 0 & 1 \\ 0 & 1 \\ 0 & 1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \end{bmatrix} = \begin{bmatrix} 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ -1 & -1 & -1 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 \\ -1 & -1 & -1 & 0 & 1 \end{bmatrix} \begin{bmatrix} \theta_A \\ \theta_B \\ \theta_C \\ \beta_1 \\ \beta_2 \end{bmatrix}$$

Quantitative outcome models

Impose a linear constraint on the θ_j through $\sum_{j=1}^J \theta_j = 0$.

Can now perform least-squares regression to obtain $\hat{\boldsymbol{\theta}}_{-1}$ and $\hat{\boldsymbol{\beta}}$.

Once $\hat{\boldsymbol{\theta}}_{-1}$ is determined, can then compute $\hat{\boldsymbol{\theta}} = \mathbf{W}\hat{\boldsymbol{\theta}}_{-1}$. See the normal model lecture notes.

	Normal.ability.est	Normal.ability.stder
1	-0.1875	0.6123724
2	0.6125	0.6123724
3	0.2625	0.6123724
4	-0.6875	0.6123724

Quantitative outcome models

Perform ridge regularization on the θ_j (but not on the β_i).

This basically involves finding the maximum of the penalized log-likelihood, that is

$$\log \tilde{\mathcal{L}}(\boldsymbol{\theta}, \boldsymbol{\beta}, \sigma^2 | \mathbf{X}, \mathbf{y}, \lambda) = \log \mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\beta}, \sigma^2 | \mathbf{X}, \mathbf{y}) - \lambda \sum_{j=1}^J \theta_j^2$$

where, up to an additive constant,

$$\begin{aligned} & \log \mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\beta}, \sigma^2 | \mathbf{X}, \mathbf{y}) \\ &= -\frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}_\theta \boldsymbol{\theta} - \mathbf{X}_\beta \boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}_\theta \boldsymbol{\theta} - \mathbf{X}_\beta \boldsymbol{\beta}) \end{aligned}$$

Example

7 x 1 sparse Matrix of class "dgCMatrix"

```
              s0
(Intercept)  .
V1           3.350
V2           3.325
V3          -0.178
V4           0.582
V5           0.250
V6          -0.654
```

GLMNET shrinks the coefficients a little bit when applying a ridge penalty.

Run ridge regression using `glmnet`.

- Use cross-validation to estimate the penalty/tuning parameter, λ .
- Set `intercept=F`.
- Make sure to prevent the coefficients to the game indicators from being shrunk (use the `penalty.factor` argument)

	LM.est	GLMNET.est
1	-0.1875	-0.1783
2	0.6125	0.5825
3	0.2625	0.2496
4	-0.6875	-0.6538

Quantitative outcome models: Comments

- Attractive features of this model:
 - Doesn't require every competitor to participate in every game.
 - No obvious bias if certain competitors choose which games in which to participate.
- Main concern: Normality assumption
 - Should perform residual diagnostics to assess the normality assumption, e.g., residual plots, QQ-plot of residuals to check for normality, etc.
 - Can always try transforming the scores (e.g., log transformation) to improve normality of the residuals.
 - Modern approaches can find a transformation automatically. See Che and Glickman (2024) for such an approach (applied to time-varying abilities).
 - If scores are suspected to follow some other distribution, the approach above can be adapted in a straightforward manner using generalized linear modeling.

Example: PGA Tour golf game scores



Example: PGA Tour golf game scores



- Data set includes 15 active golfers who competed in 2019 PGA Tour golf tournaments.
- Recorded performances across 34 golf tournaments in 2019.
- Each golfer participated in 17 to 21 of the events, so fairly balanced participation.
- The score for each event is the total strokes over par, i.e., the number of strokes above what is considered expected for that golf course. Multiplied this by -1 in our analyses so that positive values indicate better results.


Example: PGA Tour golf game scores




Rory Sabbatini
Lucas Glover
Vaughn Taylor
Brandt Snedeker
Webb Simpson
Billy Horschel
Kevin Kisner
Jason Kokrak
Nick Taylor
Patrick Reed
Hideki Matsuyama
Adam Schenk
J.T. Poston
Joaquin Niemann
Sungjae Im

Golfers in the
data set

Histogram of residuals from normal model




Normal Q-Q Plot



Example: PGA Tour golf game scores

	Normal.ability.est	Normal.ability.stder
Webb Simpson	2.882	0.933
Hideki Matsuyama	2.874	0.929
Lucas Glover	0.700	0.879
Sungjae Im	0.605	0.833
Billy Horschel	0.420	0.929
Patrick Reed	0.324	0.936
Jason Kokrak	0.320	0.869
Rory Sabbatini	0.260	0.924
Kevin Kisner	0.010	0.927
Brandt Snedeker	-0.469	0.879
Vaughn Taylor	-1.221	0.904
Adam Schenk	-1.282	0.937
J.T. Poston	-1.434	0.901
Joaquin Niemann	-1.857	0.879
Nick Taylor	-2.131	0.888

Example: PGA Tour golf game scores



Estimates with 1-SE
intervals

Ranking models

The situation: Same as before, but with one tweak

- Suppose we have J competitors, a subset of whom compete in n multiplayer games.
- For example, suppose we observe J athletes who compete in n races, and that we measure the finish time for each race.
- For game $i = 1, \dots, n$, we do not observe any score for player j in game i , or at least we ignore any score measured. Instead, we only observe the rank order for players in game i .

Ranking models: Why ignore scores if they exist?

- Clearly recording scores in a multicompetitor game (e.g., race times) is providing more information than just recording rank orderings.
- Stands to reason that modeling the scores is going to result in more accurate and more precise inferences about player strength, and more accurate predictions.
- However, the key issue in making accurate inferences about strength is that the score model is reliable!
 - If our choice of a model, e.g., normal distribution for scores, is inappropriate, inferences about player strengths may not be reliable.
 - Using rank orderings instead is a robust approach (robust to the choice of a score model). However, there may be some degradation in inferences relative to models for score outcomes.

Ranking models: Motivating example

Suppose four players, A , B , C and D , compete in two races.

The results were:

Race 1: $B \succ C \succ A \succ D$

Race 2: $B \succ A \succ C \succ D$

These are just the rank orderings from the earlier stylized example, but ignoring the scores.

Ranking models: Motivating example

Race 1: $B \succ C \succ A \succ D$

Race 2: $B \succ A \succ C \succ D$

Would like to assume player abilities $\theta_A, \theta_B, \theta_C, \theta_D$ that can be estimated from the race outcomes.

But how?

Ranking models: Motivating example

Idea:

Decompose a single multicompetitor event into all pairwise comparisons, and then use methods discussed previously (linear paired comparison models, such as Bradley-Terry) to estimate unknown strength parameters.

Race 1: $B \succ C \succ A \succ D$

Race 2: $B \succ A \succ C \succ D$

would become

Race 1: $(B \succ C), (B \succ A), (B \succ D), (C \succ A), (C \succ D), (A \succ D)$

Race 2: $(B \succ A), (B \succ C), (B \succ D), (A \succ C), (A \succ D), (C \succ D)$

Ranking models: Motivating example

Problem with this approach:

- Acts as though one multicompetitor game was equivalent to many head-to-head games.
- Analyzing the data this way may give accurate estimates of ability, but the uncertainty of the estimates will be grossly underestimated (that is, the estimates would appear overly precise).
- Furthermore, it is not obvious how to make probability predictions for future rankings.
- Finally, using a paired comparison approach is not necessary, as we shall soon see.

Ranking models: A generative approach

Suppose players A, B, C and D are about to compete in game i .

Assume (unknown) player abilities $\theta_A, \theta_B, \theta_C, \theta_D$.

Key idea: For game i , let y_{iA}, y_{iB}, y_{iC} , and y_{iD} be the latent numerical performances by the four players, where, independently,

$$y_{iA} \sim F(\theta_A)$$

$$y_{iB} \sim F(\theta_B)$$

$$y_{iC} \sim F(\theta_C)$$

$$y_{iD} \sim F(\theta_D)$$

for some distribution F that depends on the given parameter.

Ranking models: A generative approach

Suppose that $y_{iB} > y_{iC} > y_{iA} > y_{iD}$. Then the observed outcome would be the rank ordering

$$B \succ C \succ A \succ D$$

Thus, if we knew the strength parameters, then to generate a particular rank ordering we need to

1. simulate latent numerical scores, and then
2. rank order the scores.

Of course, we never see the latent numerical scores.

This approach to modeling rank orderings is sometimes called the Thurstonian model for rankings.

Ranking models: A generative approach

The probability of the rank ordering $B \succ C \succ A \succ D$, given the strength parameters, is given by

$$\Pr(B \succ C \succ A \succ D) = \Pr(y_{iB} > y_{iC} > y_{iA} > y_{iD} | \theta_A, \theta_B, \theta_C, \theta_D) = \\ \int \int \int \int_{y_{iB} > y_{iC} > y_{iA} > y_{iD}} f(y_A | \theta_A) f(y_B | \theta_B) f(y_C | \theta_C) f(y_D | \theta_D) dy_A dy_B dy_C dy_D$$

where the f 's are the densities for the latent performance distributions.

Is this multiple integral feasible to calculate?

Plackett-Luce model

Let the density for the latent performance distribution be

$$f(y|\theta) = \exp\left(-\left(y - \theta + e^{-(y-\theta)}\right)\right)$$

that is, the Gumbel distribution with location parameter θ .

In other words, for our example we assume, independently,

$$\begin{aligned} f(y_A|\theta_A) &= \exp\left(-\left(y_A - \theta_A + e^{-(y_A-\theta_A)}\right)\right) \\ &\quad \vdots \\ f(y_D|\theta_D) &= \exp\left(-\left(y_D - \theta_D + e^{-(y_D-\theta_D)}\right)\right) \end{aligned}$$

Plackett-Luce model

This distributional assumption leads to

$$\Pr(B \succ C \succ A \succ D) = \Pr(y_{iB} > y_{iC} > y_{iA} > y_{iD} | \theta_A, \theta_B, \theta_C, \theta_D) =$$

$$\left(\frac{e^{\theta_B}}{e^{\theta_B} + e^{\theta_C} + e^{\theta_A} + e^{\theta_D}} \right) \left(\frac{e^{\theta_C}}{e^{\theta_C} + e^{\theta_A} + e^{\theta_D}} \right) \left(\frac{e^{\theta_A}}{e^{\theta_A} + e^{\theta_D}} \right) =$$

$$\Pr(B \text{ outperforms } C, A, D) \Pr(C \text{ outperforms } A, D) \Pr(A \text{ outperforms } D)$$

Important to note that this result relies on the non-trivial result that

$$\Pr(C \text{ outperforms } A, D) = \Pr(C \text{ outperforms } A, D | \text{B performs best})$$

$$\Pr(A \text{ outperforms } D) = \Pr(A \text{ outperforms } D | \text{B, C outperform A and D})$$

Plackett-Luce model

In general, consider J players with strength parameters $\theta_1, \theta_2, \dots, \theta_J$.

Let $\{j_1, j_2, \dots, j_M\} \subset \{1, 2, \dots, J\}$ be a permutation of the set of M players where player j_1 comes in first place, player j_2 comes in second place, etc.

Then the Plackett-Luce model (also called the rank-order logit model) asserts

$$\Pr(j_1 \succ j_2 \succ \dots \succ j_M) = \prod_{m=1}^M \frac{e^{\theta_{j_m}}}{\sum_{\ell=m}^M e^{\theta_{j_\ell}}}$$

When $M = 2$, the model becomes

$$\Pr(j_1 \succ j_2) = \prod_{m=1}^2 \frac{e^{\theta_{j_m}}}{\sum_{\ell=m}^2 e^{\theta_{j_\ell}}} = \frac{e^{\theta_{j_1}}}{e^{\theta_{j_1}} + e^{\theta_{j_2}}}$$

which is the Bradley-Terry model! So the BT model is a special case of PL.

Plackett-Luce model: Comments

- The Plackett-Luce model satisfies Luce's "choice" axiom. This says that the relative rankings of a subset of players is independent of other players not in the rankings. That is, for example,

$$\Pr(B \succ C) = \Pr(B \succ C \mid A, D \text{ outperform } B, C)$$

- In fact, Yellott (1977) showed that the Gumbel distribution is the only distribution in a Thurstonian model for rankings that satisfies Luce's choice axiom.
- Gumbel distribution assumption also has the advantage of corresponding to probabilities of ranking orderings that can be written down in closed form.
- Assuming other distributions for performance distributions, e.g., normal, is much more complicated and does not correspond to closed-form expressions for the probability of a rank ordering.

Plackett-Luce model: Covariates

The Plackett-Luce model can be easily adapted to include covariates.

Let \mathbf{x}_{j_m} be a vector of (endogenous) covariates for competitor j_m and let \mathbf{u} be a vector of (exogenous) game-specific covariates. Then letting $\boldsymbol{\beta}$ be the associated vector of endogenous covariate effects, and letting $\boldsymbol{\gamma}_{j_m}$ be a vector of exogenous covariate effects, the modified Plackett-Luce model is given by

$$\Pr(j_1 \succ j_2 \succ \cdots \succ j_M) = \prod_{m=1}^M \frac{e^{\theta_{j_m} + \mathbf{x}_{j_m}^T \boldsymbol{\beta} + \mathbf{u}^T \boldsymbol{\gamma}_{j_m}}}{\sum_{\ell=m}^M e^{\theta_{j_\ell} + \mathbf{x}_{j_\ell}^T \boldsymbol{\beta} + \mathbf{u}^T \boldsymbol{\gamma}_{j_\ell}}}$$

As with paired comparison models, it is important to distinguish between endogenous and exogenous covariates when estimating strength.

Plackett-Luce model: Inference

For game i , the contribution to the likelihood is

$$\Pr(j_{i1} \succ j_{i2} \succ \cdots \succ j_{iM_i}) = \prod_{m=1}^{M_i} \frac{e^{\theta_{ij_m}}}{\sum_{\ell=m}^M e^{\theta_{ij_\ell}}}$$

so that the log-likelihood is

$$\begin{aligned} \log \mathcal{L}(\boldsymbol{\theta}|\mathbf{y}) &= \sum_{i=1}^n \log \Pr(j_{i1} \succ j_{i2} \succ \cdots \succ j_{iM_i}) = \\ &= \sum_{i=1}^n \sum_{m=1}^{M_i} \left(\theta_{ij_m} - \log \sum_{\ell=m}^M e^{\theta_{ij_\ell}} \right) \end{aligned}$$

Plackett-Luce model: Inference

- Can perform maximum likelihood, but need to impose the usual linear constraint on the θ_j because the θ_j are unidentifiable otherwise.
- Can also perform ridge-penalized maximum likelihood. This involves maximizing


$$\log \tilde{\mathcal{L}}(\boldsymbol{\theta}|\mathbf{y}, \lambda) = \sum_{i=1}^n \sum_{m=1}^{M_i} \left(\theta_{ij_m} - \log \sum_{\ell=m}^M e^{\theta_{ij_\ell}} \right) - \lambda \sum_{j=1}^J \theta_j^2$$

- Can perform Bayesian inference with independent normal prior densities on the θ_j , that is $\theta_j \sim N(0, \sigma^2)$.

Example: PGA Tour golf game scores

	PlackettLuce.est	PlackettLuce.stderr
Sungjae Im	-2.476	0.211
Jason Kokrak	-2.570	0.223
Hideki Matsuyama	-2.635	0.234
Lucas Glover	-2.654	0.226
Brandt Snedeker	-2.668	0.224
Rory Sabbatini	-2.716	0.230
Vaughn Taylor	-2.748	0.229
Webb Simpson	-2.770	0.238
Billy Horschel	-2.797	0.240
Joaquin Niemann	-2.806	0.224
J.T. Poston	-2.810	0.232
Kevin Kisner	-2.813	0.239
Patrick Reed	-2.823	0.237
Nick Taylor	-2.907	0.222
Adam Schenk	-2.923	0.238


Example: PGA Tour golf game scores



Estimates with 1-SE
intervals

Overlap of intervals for
Plackett-Luce model
compared to normal model.

Example: PGA Tour golf game scores



Relationship between normal estimates of ability and Plackett-Luce estimates.

Correlation of estimates is 0.482. Not surprising given the uncertainty in the Plackett-Luce estimates.

References:

Glickman, M. E., & Hennessy, J. (2015). A stochastic rank ordered logit model for rating multi-competitor games and sports. *Journal of Quantitative Analysis in Sports*, 11(3), 131-144.

Luce, R. D. (1959). *Individual Choice Behavior: A Theoretical Analysis*. New York: Wiley.

Plackett, R. L. (1975). "The Analysis of Permutations." *Appl. Statist* 24 (2): 193–202.

Yellott Jr, J. I. (1977). The relationship between Luce's choice axiom, Thurstone's theory of comparative judgment, and the double exponential distribution. *Journal of Mathematical Psychology*, 15(2), 109-144.